

Neural Network Architectures

Putting everything together

Yordan Darakchiev

Technical Trainer

iordan93@gmail.com





sli.do

#DeepLearning

Table of Contents


- Architectures: review
- Transfer learning
- Semi-supervised methods
- Image captioning

Architectures

A recap on the popular ones

VGG-19


- Input shape: $[n_{img}, 224, 224, 3]$
- Output shape: $[n_{img}, 1000]$
- Total params: 143 667 240
- Input flow
 - Read image
 - Resize
 - Preprocess
- Predict
- Decode predictions



```
from tensorflow.keras.applications.vgg19 import \
    preprocess_input, VGG19, decode_predictions
```

Inception v3

- Input shape: $[n_{img}, ?, ?, 3]$
 - Originally $? = 299$
- Output shape: $[n_{img}, 1000]$
- Total params: 23 851 784

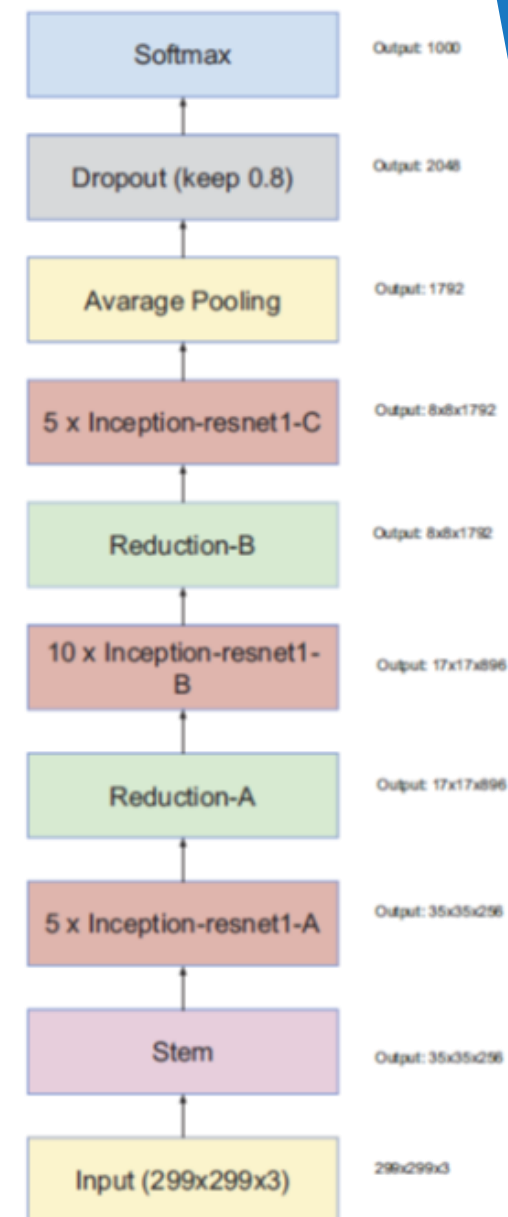


```
from tensorflow.keras.applications.inception_v3 import \
InceptionV3, decode_predictions, preprocess_input
```

Inception-ResNet v2


- Input shape: $[n_{img}, ?, ?, 3]$
 - Originally $? = 299$
 - We don't need to resize the image, just apply preprocessing
- Output shape: $[n_{img}, 1000]$
- Total params: 55 873 736

```
from tensorflow.keras.applications.inception_resnet_v2 \
import preprocess_input, InceptionResNetV2, decode_predictions
```



Object Localization

- Input: image; output: bounding box (x, y, w, h)
 - Regression
- Classification and localization
 - Simplest case: 1 object
 - Output a vector: $[p, x, y, w, h, c_1, c_2, \dots, c_k]$
 - $p = 0 \Rightarrow$ no object detected; we don't care about the other numbers
 - $p = 1 \Rightarrow$ object detected; class: c_1, \dots, c_k ; bounding box x, y, w, h
 - Metrics: usually [IoU](#) (or Euclidean distance)
- Implementations: [YOLO](#) (You Only Look Once)
 - Also: [R-CNN](#) (Region-proposing network)



Transfer Learning


- For a new problem similar to another one that we have already solved, we can reuse the weights
- Prerequisites
 - Small training set on our new model
 - Similar task (i.e. image description)
- Algorithm
 - Remove the last r layers
 - "Freeze" the weights of the remaining layers (`trainable = False`)
 - I.e. use them as a fixed function
 - Add one or more (r') layers
 - Retrain the model (this will update only the last r' layers)

Semi-Supervised Methods

Venturing into unsupervised land

Autoencoders

- NNs which learn to reconstruct their input
 - We care about the latent representation h (like CNNs)
- Encoder / Decoder are simply NNs (can be CNNs, can use multiple layers)
- Main advantages
 - Dimensionality reduction
 - Denoising
- Loss function: difference between x and \tilde{y} (MSE works well)
- Denoising autoencoder
 - Can be used for images, audio, text, etc.
 - Add noise to x to create x_{noise} , compare \tilde{y} to x (**not** to x_{noise})



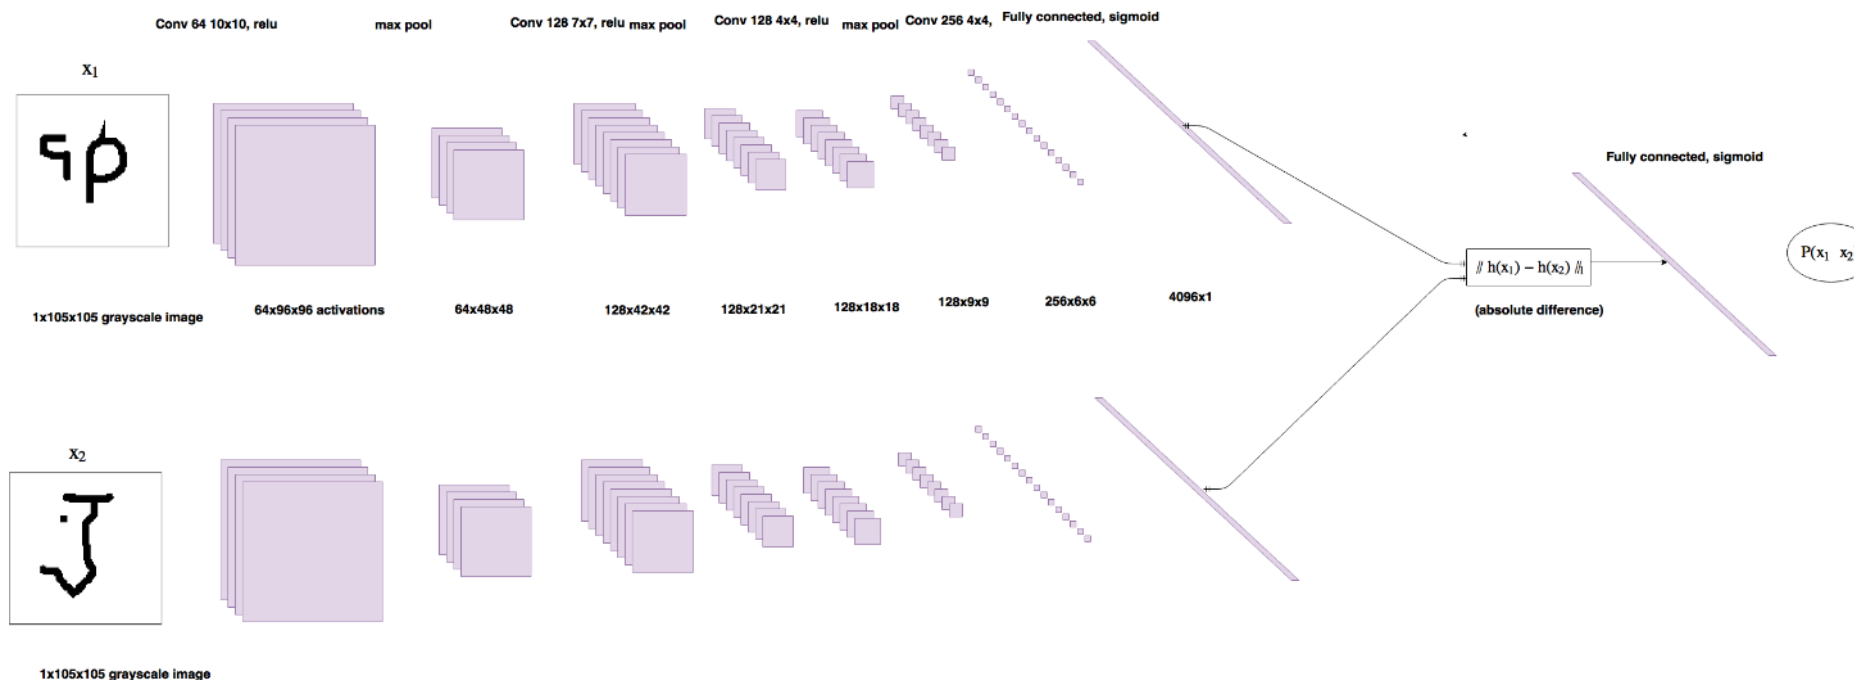
One-Shot Learning

- Example: Facial recognition
 - Input: image of a face; background data: allowed faces
 - Output: Does the image match any of the allowed faces?
- We aren't allowed to train on many pictures
 - We usually have one picture per person
- This concept generalizes to other multi-class classification tasks
- Solution: **Siamese networks**
 - Two identical networks receive two images and compute two vectors
 - The distance between the vectors is their **(dis)similarity score**
 - Dimensionality reduction technique (also similar to clustering)

One-Shot Learning (2)

- Training: **triplet loss** function

- Input images: anchor a , positive p , negative n
- Output: vectors e_a, e_p, e_n
- Loss: low distance from e_a to e_p , high distance from e_a to e_n :
$$L(a, p, n) = \max(d(e_a, e_p) - d(e_a, e_n) + m, 0)$$
 - m – margin (similar to SVMs, allows us to distinguish better)



Novelty Detection

- Similar to one-class SVM ([Chalapathy et al., 2018](#))
- Main idea
 - Assign a confidence score to samples
 - Loss function – minimize distances
- Two approaches
 - If samples have a confidence score (as output)
⇒ we can learn even if we have samples of only one class
 - If we don't have a confidence score, we can use similarity measures (i.e. similar embedding vectors)
 - Autoencoder




Image Captioning

Describing images to humans

Image Captioning

- [Vinyals et al., 2015 \(code\)](#)
- Similar to machine translation
 - Encoder: CNN instead of RNN
 - Decoder: RNN (LSTM)

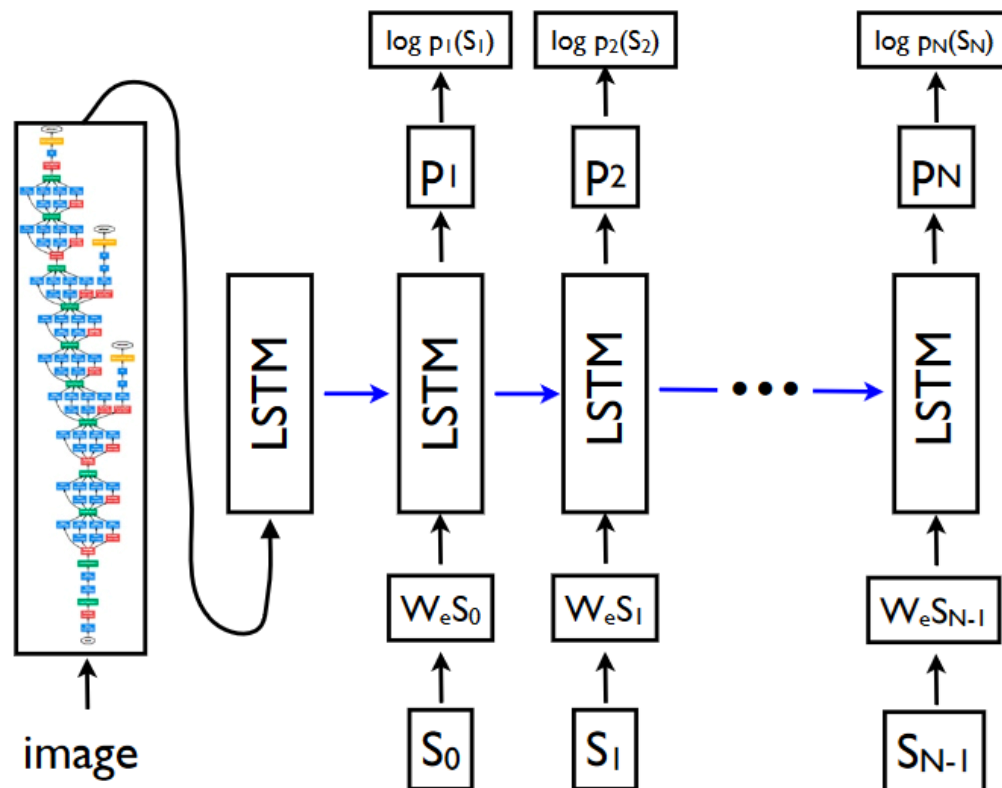




Image Captioning (2)

- Extension: [Karpathy and Li, 2015](#)
 - Goal: propose regions and describe them individually
 - R-CNN to get regions
 - Bi-directional RNN to generate sentences



- Both embeddings use the same-dimensional space



Summary

- Architectures: review
- Transfer learning
- Semi-supervised methods
- Image captioning

The image features a white background with two blue decorative bars. The top bar is a solid blue strip. The bottom bar is a gradient of blue, transitioning from a lighter shade on the left to a darker shade on the right. The word "Questions?" is centered in a blue, sans-serif font.

Questions?